

Firstname Lastname

Toronto, CA

email@example.com

+1 000 000 0000

linkedin.com/in/username

Education

University of Toronto, Bachelor of Applied Science

09/2021 - 05/2026

Machine Intelligence, cGPA: **3.94**/4.00

- Relevant Courses: Data Structures & Algorithms, Probability and Statistics, Reinforcement Learning, Digital and Computer Systems, Introduction to Machine Learning, Matrix Algebra and Optimization, Probabilistic Reasoning, Software and Neural Networks, Introduction to Databases.

Publications

- **Retrieval-Augmented Conversational Recommendation with Prompt-Based Semi-Structured Natural Language State Tracking**
First-author, *ACM SIGIR*, 2024.
- **Elaborative Subtopic Query Reformulation for Query-Driven Recommendation**
Co-author, *ACM SIGIR*, 2025 (under review).

Experience

SWE/SDE/MLE/MLops Intern, Modiface – Toronto, CA

May 2024 - Present

- Engineered a robust data processing pipeline to augment a synthetic 3D-face dataset (25M images) using stable diffusion with Python and Shell scripting, achieving a 17
- Developed a seamless model conversion tool to translate in-house computer vision models from PyTorch to TensorFlow, ensuring consistent cross-platform integration, reducing latency by over 10
- Optimized inference performance of generative AI models (including GANs and diffusion-based models) by integrating OpenVINO, reducing latency by over 40
- Developed a Python-based fine-tuning pipeline integrating existing training conversation pairs and formatting scripts to enhance beauty chatbot by replacing the in-context learning approach
- Refactored skin diagnosis repository using object-oriented design, consolidating recurring patterns into a unified sign registry and score normalization framework, enhancing code maintainability and consistency.
- Finetuned Stable Diffusion LoRA models to improve task-specific GenAI performance, improving accuracy by 40

SWE/SDE/MLE/Applied ML/ML Research intern, Data-Driven Decision Making Lab
(UofT) – Toronto, CA

April 2023 - September 2023

- Designed and implemented a RAG-based chatbot system leveraging LLM APIs for dynamic natural language interactions and personalized responses.
- Architected the application using object-oriented design patterns, creating a modular system with 10+ interchangeable components for different NLP workflows.
- Developed core infrastructure in Python with 5+ API integrations, implementing rate limiting and caching mechanisms to handle 500+ RPM.
- Built a custom data pipeline using FAISS vector database to process and embed 1M+ product entries for real-time retrieval.
- Retrieval-Augmented Conversational Recommendation System (SIGIR 2024, 1st Author): Developed an open source RAG-based conversational recommendation system powered by GPT-4.

Software Engineer Intern, Voith Hydro – Montreal, CA

May 2022 - September 2022

- Engineered a SharePoint infrastructure solution to centralize engineering resources, deployed to 200+ users across the engineering department
- Automated material specification analysis by creating Python scripts (Pandas/NumPy) to process CSV datasets, generating standardized reports that reduced manual review time.
- Conducted user research interviews with 15+ engineers to optimize UI/UX flow, resulting in 95